

Computer Lab 4

In this computer lab we will have to achieve the following tasks/learning outcomes:

- import data
- understand some important features of the dataset
- select a subset of data
- estimate a Diff-in-Diff estimator to evaluate the effectiveness of a sugar tax

The example is the example discussed in Demonstration CLass 2 and comes from the Doing Economics Project 3.

Preparing your workfile

We add the basic libraries needed for this week's work:

```
library(tidyverse)    # for almost all data handling tasks
library(ggplot2)      # to produce nice graphiscs
library(stargazer)    # to produce nice results tables
library(haven)        # to import stata file
library(AER)          # access to HS robust standard errors
library(readxl)       # enable the read_excel function
library(plm)          # enable panel data methods
```

You should also save the separately supplied `stargazer_HC.r` file in your working directory. This will make it straightforward to estimate and compare regressions with robust standard errors. Once you have done that you should include the following line into your code which basically makes this function available to you.

```
source("stargazer_HC.r") # includes the robust regression
```

We will first look at price data from the treatment group (stores in Berkeley) to see what happened to the price of sugary and non-sugary beverages after the tax.

- Download the data from the Global Food Research Program's website, and select the 'Berkeley Store Price Survey' Excel dataset. Then upload the dataset into R.
- The first tab of the Excel file contains the data dictionary. Make sure you read the data description column carefully, and check that each variable is in the Data tab.

Data Import and tidying

This is an xlsx file and we use the `readxl` package to import this file. We also install (if you haven't done so yet) and load the `tidyverse` library as this includes packages that we will use later (piping and graphing). There are two worksheets, one with some info on the variables ("Data Dictionary") and the other with the data ("Data"). It is useful to import both into R so that you have the additional data info available while you are working without having to go back into EXCEL.

Either use the R help function or google to find out how to use the `sheet =` option inside the `read_excel` function.

```
var_info <- XXXX("sps_public.xlsx", XXXX="Data Dictionary")
dat <- XXXX("sps_public.xlsx", XXXX="Data")
```

Let's look at the datatypes this import process as assigned to the respective variables.

```
str(XXXX)
```

Confirm that all variables which should be numbers, in particular prices, are indeed `num` variables. But then there are variables that may be useful to have as categorical (factor in R speak) variables. Namely `type`, `taxed`, `supp`, `store_id`, `store_type`, `type2` and `product_id`. All these variables have been expressed as numbers but really are categorical variables. So let's convert them to factor variables.

The following code converts all but the `store_type` variable. Where the data are initially `num` we supply a `label` option which gives sensible names to the different levels.

```
dat$type <- factor(dat$type)
dat$taxed <- factor(dat$taxed,
                    labels=c("not taxed", "taxed"))
dat$supp <- factor(dat$supp,
                  labels=c("Standard", "Supplemental"))
dat$store_id <- factor(dat$store_id)
dat$type2 <- factor(dat$type2)
dat$product_id <- factor(dat$product_id)
```

Before converting the `store_type` variable, let's see what information we have on this variable

```
var_info[var_info$`Variable Name` == "store_type",]
```

```
## # A tibble: 1 x 3
##   `Variable Name` Type Description
##   <chr>           <chr> <chr>
## 1 store_type      Num   Store Type: 1(Large Supermarket), 2(Small Supermarket),~
```

Let's change the `store_type` variable to a factor.

```
dat$store_type <- XXXX(dat$store_type,
                      labels=c("Large Supermarket",
                                XXXX,
                                XXXX,
                                XXXX))
```

There is another variable `time` which is characterised as a `chr` variable. This stands for characters as in letters and numbers. It may be useful to change this also into a factor variable. Before we do this we will check which values this variable takes by using the `unique` command.

```
unique(dat$time)
```

```
## [1] "DEC2014" "JUN2015" "MAR2015"
```

If you read the Silver et al. (2016) paper, and in particular if you have looked at the timeline, then you will have realised that the third survey wasn't in March 2015 but in March 2016. So the data have been labelled incorrectly. We shall therefore change all the values "MAR2015" to "MAR1016".

```
dat$time[XXXX == "MAR2015"] <- XXXX
```

We can now change the `time` variable into a factor variable.

```
dat$time <- factor(dat$time)
```

Understanding the data

How many different stores and how many different products have been included in the survey? There are different ways to achieve this. One way is to use the `unique` function in combination with the `length` function.

Confirm the following:

```
## [1] "Number of Stores: 26"
## [1] "Number of Products: 247"
```

Execute the following command and interpret its outcome. You may want to use the help function or google to understand what the `tally()` function does in this pipe.

```
table1 <- dat %>% group_by(store_type,taxed,time) %>%
  tally() %>%
  spread(time,n) %>%
  print(n=Inf)
```

Now use the `tally` function to produce a frequency table showing the number of each product type (`type`), with product type (`type`) as the row variable and time period (`time`) as the column variable.

```
table2 <- XXXX %>% group_by(XXXX,XXXX) %>%
  tally() %>%
  spread(XXXX,n) %>%
  print(n=Inf)
```

```
## # A tibble: 13 x 4
## # Groups:   type [13]
##   type      DEC2014 JUN2015 MAR2016
##   <fct>      <int>    <int>    <int>
## 1 ENERGY         56         58         49
## 2 ENERGY-DIET    49         54         35
## 3 JUICE           70         64         52
## 4 JUICE DRINK     19         17          6
## 5 MILK            63         61         53
## 6 SODA           239        262        215
## 7 SODA-DIET      128        174        127
## 8 SPORT          11         16         12
## 9 SPORT-DIET      2          2          NA
## 10 TEA            52         45         41
## 11 TEA-DIET        6          6          8
## 12 WATER          48         38         34
## 13 WATER-SWEET     1          1          1
```

Can you confirm that we have, altogether 177 observations for milk?

To explore the data a bit further let's look at all observations for the "TEA-DIET" category. In particular we want to see the `product_id`, `store_id` and `time` for these observations and only those from either "DEC2014" or "MAR2016" (should be 14). It would also be best to sort the data according to `store_id`.

```
dat %>% XXXX(type == XXXX & time %in% c(XXXX, XXXX)) %>%
  select(store_id, product_id, time) %>%
  arrange(XXXX) %>%
  print()
```

For how many stores do you have price observations in that product category for both these periods? Only one (`store_id`, 14 and 24). And only in store 24 did we record the price for the same `product_id`.

In order to establish whether prices change following the tax introduction we want to restrict our data to only those products (identified by `product_id`) for which we have observations for all three periods (DEC2014, JUN2015 and MAR2016) it is best to create a new variable `period_test` which takes the value 1 (or TRUE) if we have observations for all periods and 0 (FALSE) otherwise for a product in a particular store. We call variables like this boolean variables.

It turns out that this is not a task that is straightforward and the easiest way to achieve this is with a loop. Execute the code below. It is not crucial to understand everything now, but do, later, take some time to understand what is going on.

```
dat$period_test <- NA

sid_list = unique(dat$store_id) # list of all store ids
pid_list = unique(dat$product_id) # list of all product ids

for (s in sid_list) {
  for (p in pid_list) {
    temp <- subset(dat, product_id == p & store_id == s)
    temp_time <- temp$time
    # test will take value TRUE if we have obs from all three periods, FALSE otherwise
    test <- (any(temp_time == "DEC2014") & any(temp_time == "JUN2015") & any(temp_time == "MAR2016"))

    dat$period_test[dat$product_id == p & dat$store_id == s] <- test
  }
}
```

Now we are in a position to remove all products that have not been observed in all three periods. We define a new data frame `dat_c` with only products that are observed in all three periods in a particular store. Also, we will only select observations for which the variable `supp` takes the value “Standard”. This is what we then call a balanced panel.

```
dat_c <- dat %>% XXXX(period_test == XXXX & supp == XXXX)
```

You should confirm that you are left with 939 observations.

Now we can calculate the price averages (for `price_per_oz_c`, which is the price per ounce in cents) by grouping the data according to `store_type`, `taxed` and `time`. One way to achieve this is to use piping:

```
table_res <- dat_c %>%
  group_by(taxed, store_type, time) %>%
  summarise(n = length(price_per_oz_c), avg.price = mean(price_per_oz_c)) %>%
  spread(time, avg.price) %>%
  print()
```

```
## # A tibble: 8 x 6
## # Groups:   taxed, store_type [8]
##   taxed    store_type      n DEC2014 JUN2015 MAR2016
##   <fct>    <fct>    <int>   <dbl>   <dbl>   <dbl>
## 1 not taxed Large Supermarket    36    11.2    11.5    11.7
## 2 not taxed Small Supermarket    70    13.7    13.8    13.4
## 3 not taxed Pharmacy             18    15.2    16.1    15.4
## 4 not taxed Gas Station           12    16.9    17.0    17.0
## 5 taxed    Large Supermarket    36    15.6    16.9    16.7
## 6 taxed    Small Supermarket   101    15.9    16.0    15.5
## 7 taxed    Pharmacy             18    18.2    19.1    18.6
## 8 taxed    Gas Station           22    19.4    20.3    19.2
```

Calculating the DiD estimator

We now want to treat these as a panel dataset. You may remember that we need to define a panel dataset and let R know what variables identifies an “individual” and what variable defines the time dimension. Here an “individual” is a “particular product in a particular store”. There is no variable which can serve as such an index, but we can create it by concatenating the `store_id` and `product_id` variables.

```
dat_c$sp_id <- paste0(dat_c$store_id,"_",dat_c$product_id)
```

Have a look at values of the new variable to understand what happened. We will also have to create the interaction variable between `time` and `taxed`.

```
dat_c$did <- (dat_c$time == "MAR2016") * (dat_c$taxed == "taxed")
```

This variable will take the value of 1 for all observations which are from “MAR2016” and for taxed products. All other observations will have a value of 0.

Let’s restrict the analysis to large supermarket data and the two time periods “DEC2014” and “MAR2016”.

```
dat_c_ls <- XXXX %>% XXXX(XXXX == XXXX & XXXX %in% XXXX)
```

You should be left with 144 observations.

Then we can use this new variable `sp_id` together with the `time` variable in the definition of the panel data frame:

```
p_dat_c_ls <- pdata.frame(dat_c_ls, index = c("sp_id","time"))
```

Now we specify the regression model from which we estimate the DiD estimator.

$$y_{it} = \beta_0 + \delta d_t + \tau T_{it} + a_i + u_{it} \quad (1)$$

where y_{it} is `prize_per_oz_c`, d_t is `time`, T_{it} is `did` and a_i is our indicator variable for the individual (store-product combination), `sp_id`. This is estimated in “first-differenced” form (`model = "fd"`) using the panel estimation function `plm`.

```
mod1 <- plm(XXXX ~ XXXX + XXXX + XXXX, data = XXXX, model = "fd")
stargazer_HC(mod1)
```

```
##
## =====
##                      Dependent variable:
##                      -----
##                      price_per_oz_c
## -----
## did                      0.557
##                      (0.450)
##
## Constant                  0.510
##                      (0.318)
##
## -----
## Observations              72
## R2                        0.021
## Adjusted R2               0.007
## F Statistic               1.532 (df = 1; 70)
## =====
## Note:                      *p<0.1; **p<0.05; ***p<0.01
```

Robust standard errors in parenthesis

You can see that the estimated coefficient is 0.557, which implies that taxed products have, on average, seen price increases 0.6p/ounce higher than non-taxed products. But this difference is not statistically significantly different.

Can you find the equivalent result in the Silver et al. paper?