

R-work for Online Assessment, 2021/22

Instructions

You should work through the code below and complete it. Keep the completed code and all the resulting output. Next you should answer the questions in the online quiz. Every student will see a slightly different collection of questions (as we will randomly draw 10 questions from a pool of about 20 questions).

The questions are of four types.

- 1) Questions that merely ask you to report output from your analysis.
- 2) Some questions will ask you about R code. For example, you will see a lot of gaps (XXXX) in the code and questions may ask you how to complete the code to make the code work. Sometimes the XXXX will represent one word and on other occasions it will represent a full line (or two) of code. Other questions may ask you about the output to be produced by a particular bit of code.
- 3) The third type of questions will test your understanding of econometric issues. For example: “What is the meaning of an estimated coefficient?” “Is a particular coefficient statistically significant?”
- 4) The fourth type of question, if asked, will be on general programming issues. For example: what is the meaning of a particular error message, or, how would you search for a particular piece of information.

Preparing your workfile

Set the working directory

```
setwd("YOUR/WORKING/DIRECTORY")
```

```
setwd("C:/Rcode/RforQM/Regression")
```

We add the basic libraries needed for this week's work:

```
library(tidyverse)      # for almost all data handling tasks
library(ggplot2)        # to produce nice graphs
library(stargazer)      # to produce nice results tables
library(haven)          # to import stata file
library(AER)            # access to HS robust standard errors
library(countrycode)

source("stargazer_HC.r") # includes the robust regression display
```

Introduction

The data are an extract from the OECD.

Task 1: Data Upload - and understanding data structure

Upload the data, which are saved in the two csv files `OECD_RoadAccidents.csv` (linked from here) and `OECD_Passenger_Transport.csv` (linked from here).

```

data_acc <- XXXX("OECD_RoadAccidents.csv")
data_pt <- XXXX("OECD_Passenger_Transport.csv")

names(data_acc)
names(data_pt)

data_acc <- read_csv("OECD_RoadAccidents.csv")
data_pt <- read_csv("OECD_Passenger_Transport.csv")
names(data_acc)

## [1] "LOCATION" "INDICATOR" "SUBJECT" "MEASURE" "FREQUENCY" "TIME"
## [7] "Value"

names(data_pt)

## [1] "LOCATION" "INDICATOR" "SUBJECT" "MEASURE" "FREQUENCY" "TIME"
## [7] "Value"

```

Look at the spreadsheets and understand the data structure.

Task 2: Cleaning Accident data

Run the following lines and use the information to understand the data structure.

```

data_acc[1,]

## # A tibble: 1 x 7
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value
##   <chr>      <chr>      <chr>   <chr>   <chr>      <dbl> <dbl>
## 1 SVN      ROADACCID DEATH    NBR      A          1970   620

unique(data_acc$MEASURE)

## [1] "NBR"          "1000000HAB" "1000000VEH"

unique(data_acc$SUBJECT)

## [1] "DEATH"          "INJURE"          "ACCIDENTCASUAL"

```

In order to understand what these different measures and subjects represent you should consult the website linked above from which the data were downloaded.

Now find all the data which relate to Germany (country code: DEU) and Poland (country code: POL) in 2017.

```

data_acc %>% filter(LOCATION %in% c("DEU", "POL"), TIME == "2017")

## # A tibble: 10 x 7
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value
##   <chr>      <chr>      <chr>   <chr>   <chr>      <dbl> <dbl>
## 1 POL      ROADACCID INJURE    NBR      A          2017   39466
## 2 POL      ROADACCID ACCIDENTCASUAL NBR      A          2017   32760
## 3 POL      ROADACCID DEATH    1000000HAB A          2017    74.5
## 4 DEU      ROADACCID INJURE    NBR      A          2017   390312
## 5 DEU      ROADACCID DEATH    1000000HAB A          2017    38.5
## 6 DEU      ROADACCID ACCIDENTCASUAL NBR      A          2017   302656
## 7 POL      ROADACCID DEATH    NBR      A          2017    2831
## 8 DEU      ROADACCID DEATH    NBR      A          2017    3180
## 9 DEU      ROADACCID DEATH    1000000VEH A          2017    57.2

```

```
## 10 POL      ROADACCID DEATH      1000000VEH A      2017      95.5
```

You should see 10 rows of data.

We will only keep death data and only those data which measure the number of deaths by 1000000 inhabitants.

```
data_acc <- data_acc %>% filter(SUBJECT == "DEATH") %>%
  filter(MEASURE == "1000000HAB")
```

The only data left in the Value column is the deaths per 1,000,000 inhabitants variable. We therefore can now change the name of that variable to Deaths_p1M.

```
names(XXXX)[names(data_acc)=="XXXX"] <- "Deaths_p1M"
```

```
names(data_acc)[names(data_acc)=="Value"] <- "Deaths_p1M"
```

The datafile contains a few variables which are now redundant or we don't need any longer. We will only keep the variables we need, LOCATION, TIME and Deaths_p1M. The code below is faulty and you need to fix it

```
data_acc <- data_acc %>% Select(LOCATION, time, Deaths_p1M)
```

```
data_acc <- data_acc %>% select(LOCATION, TIME, Deaths_p1M)
```

The LOCATION variable represents the three digit country code of the respective observation. Let's add proper country names to our variables. For this we use the function `countrycode` from the `countrycode` package. The available information in the datafiles is in the LOCATION variable. It has a numeric ISO-3 format (`origin = "iso3c"`). We want to create a new variable in both of our datasets which is called `country` and contains the respective full English country name. You will have to consult the help for function `countrycode` to figure out what the destination format should be.

```
data_acc$country <- XXXX(data_acc$LOCATION, origin = "iso3c",
  destination = "XXXX")
```

```
data_acc$country <- countrycode(data_acc$LOCATION, origin = "iso3c",
  destination = "country.name")
```

Task 3: Cleaning Personal Transport data

Run the following lines and use the information to understand the data structure.

```
data_pt[1,]
```

```
## # A tibble: 1 x 7
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value
##   <chr>      <chr>      <chr>  <chr>  <chr>      <dbl> <dbl>
## 1 AUS      PASSTRANS RAIL    MLN_PKM A      1970 12473.
```

```
unique(data_pt$MEASURE)
```

```
## [1] "MLN_PKM"
```

```
unique(data_pt$SUBJECT)
```

```
## [1] "RAIL" "ROAD" "INLAND"
```

What does MLN_PKM stand for? What do the different values in SUBJECT represent? You may need to refer to the website (see link above) to check. Note that the miles traveled are given as total distance. The information is not standardised by population size. We will do this later.

You may want to check your understanding by looking at the data for Germany and Poland in 2017.

```
data_pt %>% filter(LOCATION %in% c("DEU", "POL"), TIME == "2017")
```

```
## # A tibble: 6 x 7
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value
##   <chr>      <chr>      <chr> <chr> <chr>      <dbl> <dbl>
## 1 DEU      PASSTRANS RAIL      MLN_PKM A      2017 95530
## 2 DEU      PASSTRANS ROAD      MLN_PKM A      2017 977430
## 3 POL      PASSTRANS RAIL      MLN_PKM A      2017 20319
## 4 POL      PASSTRANS ROAD      MLN_PKM A      2017 257610
## 5 DEU      PASSTRANS INLAND     MLN_PKM A      2017 1072960
## 6 POL      PASSTRANS INLAND     MLN_PKM A      2017 277929
```

*# You should see that, in a particular year in a particular country,
RAIL + ROAD = INLAND*

We basically have three variables here, RAIL travel kilometers (or better miles), ROAD travel and total travel (INLAND). We want to keep these three variables. To do so we use the `pivot_wider` command which is part of the tidyverse package.

```
data_pt <- data_pt %>% pivot_wider(names_from = SUBJECT,
                                   values_from = Value)
```

This is a useful function and you may want to remember that this functionality exists for your later work.

Display again the data for Germany and Poland in 2017 to see the difference of the datafile's setup after this operation. You should see the following:

```
data_pt %>% filter(LOCATION %in% c("DEU", "POL"), TIME == "2017")
```

```
## # A tibble: 2 x 8
##   LOCATION INDICATOR MEASURE FREQUENCY TIME RAIL ROAD INLAND
##   <chr>      <chr>      <chr> <chr>      <dbl> <dbl> <dbl> <dbl>
## 1 DEU      PASSTRANS MLN_PKM A      2017 95530 977430 1072960
## 2 POL      PASSTRANS MLN_PKM A      2017 20319 257610 277929
```

*# You should see that, in a particular year in a particular country,
RAIL + ROAD = INLAND*

The datafile contains a few variables which are now redundant or we don't need any longer. We will only keep the variables we need, LOCATION, TIME, RAIL, ROAD and INLAND.

```
data_pt <- XXXX %>% XXXX(LOCATION, TIME, RAIL, ROAD, INLAND)
```

```
data_pt <- data_pt %>% select(LOCATION, TIME, RAIL, ROAD, INLAND)
```

Now add the country name as above for `data_acc`.

```
data_pt$XXXX <- XXXX(data_pt$XXXX, origin = "iso3c",
                     destination = "XXXX")
```

```
data_pt$country <- countrycode(data_pt$LOCATION, origin = "iso3c",
                               destination = "country.name")
```

Task 4: Merging dataset

We will want to merge the two datasets `data_acc` and `data_pt` together. Before merging it is useful to review the variable names of both files and the number of observations in each.

```
nrow(data_acc)
```

```
## [1] 1368
```

```
names(data_acc)
```

```
## [1] "LOCATION" "TIME" "Deaths_p1M" "country"
```

```
nrow(data_pt)
```

```
## [1] 2598
```

```
names(data_pt)
```

```
## [1] "LOCATION" "TIME" "RAIL" "ROAD" "INLAND" "country"
```

The files have three variables in common, `LOCATION`, `TIME` and `country`. This is the information on the basis of which we want to match the data. So for instance we will want one row of data for Germany in 2017 and that row of data should contain `Deaths_p1M` from the `data_acc` dataset and `RAIL`, `ROAD` and `INLAND` from the `data_pt` dataset.

We use the `merge` function to achieve this. As you can see from the above information, the two datasets contain different numbers of rows. This is because not all the sources have complete information. To demonstrate that, let's look at Ukraine (UKR) and find the information for Ukraine from both datasets.

```
data_acc %>% filter(LOCATION == "UKR")
```

```
## # A tibble: 24 x 4
##   LOCATION TIME Deaths_p1M country
##   <chr>    <dbl>    <dbl> <chr>
## 1 UKR      1994      146. Ukraine
## 2 UKR      1995      146. Ukraine
## 3 UKR      1996      130. Ukraine
## 4 UKR      1997      118. Ukraine
## 5 UKR      1998      110. Ukraine
## 6 UKR      1999      106. Ukraine
## 7 UKR      2000      105. Ukraine
## 8 UKR      2001      123. Ukraine
## 9 UKR      2002      124. Ukraine
## 10 UKR     2003      150. Ukraine
## # ... with 14 more rows
```

```
data_pt %>% filter(LOCATION == "UKR")
```

```
## # A tibble: 31 x 6
##   LOCATION TIME RAIL ROAD INLAND country
##   <chr>    <dbl> <dbl> <dbl> <dbl> <chr>
## 1 UKR      1990  76038    NA    NA Ukraine
## 2 UKR      1991  70968    NA    NA Ukraine
## 3 UKR      1992  76196    NA    NA Ukraine
## 4 UKR      1993  75896    NA    NA Ukraine
## 5 UKR      1994  70882    NA    NA Ukraine
## 6 UKR      1995  63752    NA    NA Ukraine
## 7 UKR      1996  59080    NA    NA Ukraine
## 8 UKR      1997  54433    NA    NA Ukraine
## 9 UKR      1998  49938    NA    NA Ukraine
## 10 UKR     1999  47600    NA    NA Ukraine
## # ... with 21 more rows
```

You should see that the information for Ukraine starts in 1990 in the `data_pt` datafile but only in 1994 in the `data_acc` datafile. You can also see that the Ukraine only has information on rail travel, but not road travel.

You could decide, as we are merging data, to only keep information which is available from both datafiles, or to keep all information (Year-country combinations) which are available in at least one of the datasets. Here we decide that we want to do the latter, i.e. keep all information for now, even if it cannot be matched from the other file. So in the above case that means that we do want the information from 1990 to 1993 from Ukraine included. In effect that means that our new dataset, `data_merge` should have at least 2598 rows.

Which of the following commands does that? (Note that, as the information on the basis of which we match, country and year, is saved in identically named columns in both datafiles, we do not have to use the `by.x` and `by.y` options in the merge function).

```
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = TRUE)
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = FALSE)
data_merge <- merge(data_pt,data_acc, all.x = FALSE, all.y = FALSE)
data_merge <- merge(data_pt,data_acc, all.x = FALSE, all.y = TRUE)
```

```
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = TRUE)
```

Task 5: Calculate country averages

We now have multiple years of data for the countries in our dataset, `data_avg`. The next step is to calculate the averages for our four substantial variables (Deaths_p1M, RAIL, ROAD and INLAND).

```
data_avg <- XXXX %>% XXXX(country,LOCATION) %>%
  XXXX(avg_rail = mean(XXXX,na.rm = XXXX),
        avg_road = XXXX,
        avg_inland = XXXX,
        avg_deaths_p1M = XXXX)
```

```
data_avg <- data_merge %>% group_by(country,LOCATION) %>%
  summarise(avg_rail = mean(RAIL,na.rm = TRUE),
            avg_road = mean(ROAD,na.rm = TRUE),
            avg_inland = mean(INLAND,na.rm = TRUE),
            avg_deaths_p1M = mean(Deaths_p1M,na.rm = TRUE))
```

You got it right, if your result, `data_avg` has 59 rows and you can replicate the following information:

```
head(data_avg,10)
```

```
## # A tibble: 10 x 6
## # Groups:   country [10]
##   country      LOCATION avg_rail avg_road avg_inland avg_deaths_p1M
##   <chr>         <chr>      <dbl>   <dbl>   <dbl>         <dbl>
## 1 Albania      ALB         246.    6075.    6177.         96.9
## 2 Argentina    ARG        7765.   34058.   49191.        127.
## 3 Armenia      ARM          44.8    1945.    1987.         99.4
## 4 Australia    AUS       11658.  220449.  232107.        67.6
## 5 Austria      AUT        8792.   55652.   62935.         89.9
## 6 Azerbaijan   AZE        1304.   12856.   13838.         98.2
## 7 Belarus      BLR       11235.    NaN     NaN          134.
## 8 Belgium      BEL        8097.   98049.  106104.        100.
## 9 Bosnia & Herzegovina BIH         691.    NaN     NaN          89.4
## 10 Bulgaria     BGR        4748.   24748.   22303.        113.
```

As you can see, not all countries will have information on all variables.

As mentioned above, the information coming from the passenger transport file are not standardised by population size (avg_rail, avg_road and avg_inland), while avg_deaths_p1M already is. We now want to standardise the three variables coming from the passenger transport file by population size. For that we need to import a file with population information. That information (for more than 200 countries) is in CountryInfo.csv. Import that file and merge all the available country info into data_avg only keeping the 59 rows which we have in data_avg.

Note that you should match by the three letter country code which in CountryInfo.csv is labelled as countryCode. You will therefore, now have to use the by.x and by.y options in the merge function.

```
data_countries = read_csv("XXXX", na = "XXXX" )
data_avg <- merge(data_avg, XXXX,
                  by.x = "XXXX", by.y = "XXXX",
                  all.x = XXXX, all.y = XXXX)
```

```
data_countries = read_csv("CountryInfo.csv", na = "NA" )
data_avg <- merge(data_avg, data_countries,
                  by.x = "LOCATION", by.y = "countryCode",
                  all.x = TRUE, all.y = FALSE)
```

Now check the names of the variables in your datafile.

```
names(data_avg)
```

```
## [1] "LOCATION"      "country"      "avg_rail"     "avg_road"
## [5] "avg_inland"   "avg_deaths_p1M" "popData2019"  "continentExp"
## [9] "Land_Area_sqkm" "HealthExp"     "GDPpc"        "Obese_Pcent"
## [13] "Over_65s"     "Diabetis"
```

You should have 14 variables. Also confirm that you do have 59 rows of data. All the new country data are from the year 2019.

Task 6: Standardisation

We now need to calculate the standardised personal transport variables.

```
data_avg <- data_avg %>% XXXX(
  avg_rail_pp = avg_rail/popData2019 * 1000000,
  avg_road_pp = XXXX,
  avg_inland_pp = XXXX)
)
```

```
data_avg <- data_avg %>% mutate(
  avg_rail_pp = avg_rail/popData2019 * 1000000,
  avg_road_pp = avg_road/popData2019 * 1000000,
  avg_inland_pp = avg_inland/popData2019 * 1000000)
```

You get it right if you can replicate the following.

```
data_avg %>% filter(LOCATION %in% c("DEU", "POL")) %>%
  select(LOCATION, avg_rail_pp, avg_road_pp, avg_inland_pp)
```

```
## LOCATION avg_rail_pp avg_road_pp avg_inland_pp
## 1 DEU 752.3575 9082.167 9834.525
## 2 POL 843.5536 4885.718 5595.204
```

So the average number of miles traveled by rail per person in Germany is 752 miles per year. The average number of miles traveled on road per person in Poland is 4886. Convince yourself that the above calculation

(`avg_rail_pp = avg_rail/popData2019 * 1000000`) did deliver miles per person. recall that the travel distance coming from the passenger transport datafile was measured in Millions of miles. The population is the actual population. If we didn't multiply by 1,000,000 we would get a measure of average millions of miles traveled per person.

Lastly, we will rescale the `GDPpc` variable. It is currently defined in USD, but to simplify later analysis we want to express it in 1000 USD. So for instance Albania's `GDPpc` is USD 5224, but we want to express it as 5.224 [1000 USD].

```
data_avg <- data_avg %>% mutate(GDPpc = GDPpc/1000)
```

Look at the dataset to confirm that your operation was successful.

Task 7: Creating League Tables

Create a league table of the countries in your dataset where people travel most and least by rail and road. You want to display the 10 countries with the most and the 10 countries with the least average travel per person in both categories.

Here is the code for the table for the top 10 rail travel nations.

```
task_7a <- data_avg %>%
select(country, avg_rail_pp) %>% arrange(desc(avg_rail_pp))
head(task_7a,10)
```

```
##      country avg_rail_pp
## 1      Japan  2925.0936
## 2 Switzerland  1589.8252
## 3      Latvia  1371.0047
## 4      Russia  1318.4358
## 5      Belarus  1188.9556
## 6      Ukraine  1153.9672
## 7      France  1116.1391
## 8      Hungary  1075.6537
## 9      Austria   987.6995
## 10 Kazakhstan   915.7539
```

Repeat this for the Top 10 road travel nations and equally find similar tables for the nations travelling least (but non-zero amounts) by rail and road.

```
task_7b <- data_avg %>%
select(country, avg_road_pp) %>% arrange(desc(avg_road_pp)) %>% print()
task_7c <- data_avg %>%
select(country, avg_rail_pp) %>% arrange(avg_rail_pp) %>% print()
task_7d <- data_avg %>%
select(country, avg_road_pp) %>% arrange(avg_road_pp) %>% print()
```

Task 8: Estimate regression models - Version 1

We shall estimate the following three regression models (`mod1`)

$$avg_deaths_p1M = \gamma_0 + \gamma_1 GDPpc + w$$

and (`mod2`)

$$avg_deaths_p1M = \alpha_0 + \alpha_1 avg_rail_pp + v$$

and (mod3)

$$avg_deaths_p1M = \beta_0 + \beta_1 avg_road_pp + \beta_2 GDPpc + u$$

```
mod1 <- lm(XXXX ~ XXXX, data = data_avg)
mod2 <- lm(XXXX)
mod3 <- lm(XXXX)
stargazer_HC(mod1,mod2,mod3,omit.stat = "f")

mod1 <- lm(avg_deaths_p1M ~ GDPpc, data = data_avg)
mod2 <- lm(avg_deaths_p1M ~ avg_rail_pp, data = data_avg)
mod3 <- lm(avg_deaths_p1M ~ avg_road_pp+GDPpc, data = data_avg)
stargazer_HC(mod1,mod2,mod3,type_out="latex",omit.stat = "f")
```

% Table created by stargazer v.5.2.2 by Marek Hlavac, Harvard University. E-mail: hlavac at fas.harvard.edu
 % Date and time: Mon, Mar 14, 2022 - 09:52:20

Table 1:

	<i>Dependent variable:</i>		
	avg_deaths_p1M		
	(1)	(2)	(3)
avg_road_pp			0.003 (0.002)
GDPpc	-0.718*** (0.202)		-1.218*** (0.245)
avg_rail_pp		-0.002 (0.011)	
Constant	122.825*** (6.962)	103.763*** (6.241)	117.284*** (10.919)
Observations	55	53	42
R ²	0.288	0.001	0.369
Adjusted R ²	0.274	-0.019	0.336
Residual Std. Error	29.346 (df = 53)	33.680 (df = 51)	29.034 (df = 39)
<i>Note:</i>		*p<0.1; **p<0.05; ***p<0.01 Robust standard errors in parenthesis	

If you have done this correctly, you will find that that your estimated constant for `mod1` is 122.82, the estimated constant for `mod2` is 103.763 and that `mod3` is estimated with 42 observations.


Task 9: Plot the data

Now we want to plot the `GDPpc` versus the `avg_deaths_p1M` data in a scatter plot. Replicate the graph below.


You should change the Axis labels and add a title as shown below. If you google you should find the appropriate commands to do so. (Think carefully about the search terms.) Also, what does the `geom_smooth(method='lm')` line achieve? Figure that out by running the code without that line.

```
ggplot(data_avg, aes(x=XXXX,y=XXXX)) +
  geom_point(color = "blue") +
  geom_smooth(method='lm') +
  XXXX +
  XXXX +
  XXXX
```

```
ggplot(data_avg, aes(x=GDPpc,y=avg_deaths_p1M)) +
  geom_point(color = "blue") +
  geom_smooth(method='lm') +
  ggtitle("Accidents v GDPpc") +
  ylab("Accidents per Million population") +
  xlab("GDP per capita")
```



Finally another searching challenge for you. The picture below has added data labels to the above plot. Find a way to achieve this. It is not important that the graph looks exactly like the one below, but you do want to find out how to add data labels. Dr. Google is your friend!



And does anyone know what is going on in Luxembourg?

END OF INSTRUCTIONS