

R-work for Online Assessment, 2022/23

Instructions

You should work through the code below and complete it. Keep the completed code and all the resulting output. Next you should answer the questions in the online quiz. Every student will see a slightly different collection of questions (as we will randomly draw 10 questions from a pool of about 20 questions).

The questions are of four types.

- 1) Questions that merely ask you to report output from your analysis.
- 2) Some questions will ask you about R code. For example, you will see a lot of gaps (XXXX) in the code and questions may ask you how to complete the code to make the code work. Sometimes the XXXX will represent one word and on other occasions it will represent a full line (or two) of code. Other questions may ask you about the output to be produced by a particular bit of code.
- 3) The third type of questions will test your understanding of econometric issues. For example: “What is the meaning of an estimated coefficient?” “Is a particular coefficient statistically significant?”
- 4) The fourth type of question, if asked, will be on general programming issues. For example: what is the meaning of a particular error message, or, how would you search for a particular piece of information.

Preparing your workfile

Set the working directory

```
setwd("YOUR/WORKING/DIRECTORY")
```

```
setwd("C:/Rcode/RforQM/Regression")
```

We add the basic libraries needed for this week's work:

```
library(tidyverse)      # for almost all data handling tasks
library(ggplot2)        # to produce nice graphics
library(stargazer)      # to produce nice results tables
library(haven)          # to import stata file
library(AER)            # access to HS robust standard errors
library(countrycode)

source("stargazer_HC.r") # includes the robust regression display
```

Introduction

The data are an extract from the OECD.

Task 1: Data Upload - and understanding data structure

Upload the data, which are saved in the two csv files `OECD_RoadAccidents.csv` (linked from here) and `OECD_Passenger_Transport.csv` (linked from here).

```
data_acc <- XXXX("OECD_RoadAccidents.csv")
data_pt <- XXXX("OECD_Passenger_Transport.csv")

names(data_acc)
names(data_pt)

data_acc <- read_csv("OECD_RoadAccidents.csv")
data_pt <- read_csv("OECD_Passenger_Transport.csv")
names(data_acc)

## [1] "LOCATION" "INDICATOR" "SUBJECT" "MEASURE" "FREQUENCY"
## [6] "TIME" "Value" "Flag Codes"

names(data_pt)

## [1] "LOCATION" "INDICATOR" "SUBJECT" "MEASURE" "FREQUENCY"
## [6] "TIME" "Value" "Flag Codes"
```

Look at the spreadsheets and understand the data structure.

Task 2: Cleaning Accident data

Run the following lines and use the information to understand the data structure.

```
data_acc[1,]

## # A tibble: 1 x 8
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value `Flag Codes`
##   <chr>      <chr>      <chr>   <chr>   <chr>      <dbl> <dbl> <chr>
## 1 SVN      ROADACCID DEATH    NBR      A          1970   620 <NA>

unique(data_acc$MEASURE)

## [1] "NBR" "1000000HAB" "1000000VEH"

unique(data_acc$SUBJECT)

## [1] "DEATH" "INJURE" "ACCIDENTCASUAL"
```

In order to understand what these different measures and subjects represent you should consult the website linked above from which the data were downloaded.

Now find all the data which relate to France (country code: FRA) and Spain (country code: ESP) in 2017.

```
data_acc %>% filter(LOCATION %in% c("FRA", "ESP"), TIME == "2017")

## # A tibble: 10 x 8
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value Flag ~1
##   <chr>      <chr>      <chr>   <chr>   <chr>      <dbl> <dbl> <chr>
## 1 FRA      ROADACCID INJURE    NBR      A          2017  7.34e+4 <NA>
## 2 FRA      ROADACCID DEATH    1000000HAB A          2017  5.15e+0 <NA>
## 3 FRA      ROADACCID ACCIDENTCASUAL NBR      A          2017  5.86e+4 <NA>
## 4 ESP      ROADACCID DEATH    NBR      A          2017  1.83e+3 <NA>
## 5 FRA      ROADACCID DEATH    NBR      A          2017  3.45e+3 <NA>
## 6 ESP      ROADACCID DEATH    1000000HAB A          2017  3.93e+0 <NA>
## 7 ESP      ROADACCID ACCIDENTCASUAL NBR      A          2017  1.02e+5 <NA>
## 8 ESP      ROADACCID INJURE    NBR      A          2017  1.39e+5 <NA>
## 9 FRA      ROADACCID DEATH    1000000VEH A          2017  7.23e-1 <NA>
## 10 ESP     ROADACCID DEATH    1000000VEH A          2017  5.38e-1 <NA>
```

```
## # ... with abbreviated variable name 1: `Flag Codes`
```

You should see 10 rows of data.

We will only keep death data and only those data which measure the number of deaths by 1000000 inhabitants.

```
data_acc <- data_acc %>% filter(SUBJECT == "DEATH") %>%  
  filter(MEASURE == "1000000HAB")
```

The only data left in the Value column is the deaths per 1,000,000 inhabitants variable. We therefore can now change the name of that variable to Deaths_p1M.

```
names(XXXX)[names(data_acc)=="XXXX"] <- "Deaths_p1M"
```

```
names(data_acc)[names(data_acc)=="Value"] <- "Deaths_p1M"
```

The datafile contains a few variables which are now redundant or we don't need any longer. We will only keep the variables we need, LOCATION, TIME and Deaths_p1M. The code below is faulty and you need to fix it

```
data_acc <- data_acc %>% Select(LOCAION, Time, Deaths_p1M)
```

```
data_acc <- data_acc %>% select(LOCATION, TIME, Deaths_p1M)
```

The LOCATION variable represents the three digit country code of the respective observation. Let's add proper country names to our variables. For this we use the function `countrycode` from the `countrycode` package. The available information in the datafiles is in the LOCATION variable. It has a numeric ISO-3 format (`origin = "iso3c"`). We want to create a new variable in both of our datasets which is called `country` and contains the respective full English country name. You will have to consult the help for function `countrycode` to figure out what the destination format should be.

```
data_acc$country <- XXXX(data_acc$LOCATION, origin = "iso3c",  
  destination = "XXXX")
```

```
data_acc$country <- countrycode(data_acc$LOCATION, origin = "iso3c",  
  destination = "country.name")
```

Task 3: Cleaning Personal Transport data

Run the following lines and use the information to understand the data structure.

```
data_pt[1,]
```

```
## # A tibble: 1 x 8  
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value `Flag Codes`  
##   <chr>      <chr>      <chr>  <chr>  <chr>      <dbl> <dbl> <chr>  
## 1 AUS      PASSTRANS RAIL    MLN_PKM A        1970 12473. <NA>
```

```
unique(data_pt$MEASURE)
```

```
## [1] "MLN_PKM"
```

```
unique(data_pt$SUBJECT)
```

```
## [1] "RAIL" "ROAD" "INLAND"
```

What does MLN_PKM stand for? What do the different values in SUBJECT represent? You may need to refer to the website (see link above) to check. Note that the miles traveled are given as total distance. The information is not standardised by population size. We will do this later.

You may want to check your understanding by looking at the data for France and Spain in 2017.

```
data_pt %>% filter(LOCATION %in% c("FRA","ESP"), TIME == "2017")
```

```
## # A tibble: 6 x 8
##   LOCATION INDICATOR SUBJECT MEASURE FREQUENCY TIME Value `Flag Codes`
##   <chr>      <chr>      <chr>  <chr>  <chr>      <dbl>  <dbl> <chr>
## 1 FRA      PASSTRANS RAIL    MLN_PKM A      2017  110568 <NA>
## 2 FRA      PASSTRANS ROAD   MLN_PKM A      2017  873037 <NA>
## 3 ESP      PASSTRANS RAIL    MLN_PKM A      2017   27487 <NA>
## 4 ESP      PASSTRANS ROAD   MLN_PKM A      2017  363368 <NA>
## 5 FRA      PASSTRANS INLAND MLN_PKM A      2017  983605 <NA>
## 6 ESP      PASSTRANS INLAND MLN_PKM A      2017  390855 <NA>
```

*# You should see that, in a particular year in a particular country,
RAIL + ROAD = INLAND*

We basically have three variables here, RAIL travel kilometers (or better miles), ROAD travel and total travel (INLAND). We want to keep these three variables and display them each as a column variable. To do so we use the `pivot_wider` command which is part of the tidyverse package.

```
data_pt <- data_pt %>% pivot_wider(names_from = SUBJECT,
                                   values_from = Value)
```

This is a useful function and you may want to remember that this functionality exists for your later work.

Display again the data for France and Spain in 2017 to see the difference of the datafile's setup after this operation. You should see the following:

```
data_pt %>% filter(LOCATION %in% c("FRA","ESP"), TIME == "2017")
```

```
## # A tibble: 2 x 9
##   LOCATION INDICATOR MEASURE FREQUENCY TIME `Flag Codes` RAIL ROAD INLAND
##   <chr>      <chr>      <chr>  <chr>      <dbl> <chr>      <dbl> <dbl> <dbl>
## 1 FRA      PASSTRANS MLN_PKM A      2017 <NA>      110568 873037 983605
## 2 ESP      PASSTRANS MLN_PKM A      2017 <NA>      27487 363368 390855
```

*# You should see that, in a particular year in a particular country,
RAIL + ROAD = INLAND*

The datafile contains a few variables which are now redundant or we don't need any longer. We will only keep the variables we need, LOCATION, TIME, RAIL, ROAD and INLAND.

```
data_pt <- XXXX %>% XXXX(LOCATION, TIME, RAIL, ROAD, INLAND)
```

```
data_pt <- data_pt %>% select(LOCATION, TIME, RAIL, ROAD, INLAND)
```

Now add the country name as above for `data_acc`.

```
data_pt$XXXX <- XXXX(data_pt$XXXX, origin = "iso3c",
                     destination = "XXXX")
```

```
data_pt$country <- countrycode(data_pt$LOCATION, origin = "iso3c",
                               destination = "country.name")
```

After doing all this you should find that your data have the following dimensions:

- `data_acc`: 1513 obs and 4 variables
- `data_pt`: 2886 obs and 6 variables

Task 4: Merging dataset

We will want to merge the two datasets `data_acc` and `data_pt` together. Before merging it is useful to review the variable names of both files and the number of observations in each.

```
nrow(data_acc)
```

```
## [1] 1513
```

```
names(data_acc)
```

```
## [1] "LOCATION" "TIME" "Deaths_p1M" "country"
```

```
nrow(data_pt)
```

```
## [1] 2886
```

```
names(data_pt)
```

```
## [1] "LOCATION" "TIME" "RAIL" "ROAD" "INLAND" "country"
```

The files have three variables in common, `LOCATION`, `TIME` and `country`. This is the information on the basis of which we want to match the data. So for instance we will want one row of data for Germany in 2017 and that row of data should contain `Deaths_p1M` from the `data_acc` dataset and `RAIL`, `ROAD` and `INLAND` from the `data_pt` dataset.

We use the `merge` function to achieve this. As you can see from the above information, the two datasets contain different numbers of rows. This is because not all the sources have complete information. To demonstrate that, let's look at Ukraine (UKR) and find the information for Ukraine from both datasets.

```
data_acc %>% filter(LOCATION == "UKR")
```

```
## # A tibble: 24 x 4
##   LOCATION TIME Deaths_p1M country
##   <chr>    <dbl>    <dbl> <chr>
## 1 UKR      1994      14.6 Ukraine
## 2 UKR      1995      14.6 Ukraine
## 3 UKR      1996      13.0 Ukraine
## 4 UKR      1997      11.8 Ukraine
## 5 UKR      1998      11.0 Ukraine
## 6 UKR      1999      10.6 Ukraine
## 7 UKR      2000      10.5 Ukraine
## 8 UKR      2001      12.3 Ukraine
## 9 UKR      2002      12.4 Ukraine
## 10 UKR     2003      15.0 Ukraine
## # ... with 14 more rows
```

```
data_pt %>% filter(LOCATION == "UKR")
```

```
## # A tibble: 31 x 6
##   LOCATION TIME RAIL ROAD INLAND country
##   <chr>    <dbl> <dbl> <dbl> <dbl> <chr>
## 1 UKR      1990 76038   NA    NA Ukraine
## 2 UKR      1991 70968   NA    NA Ukraine
## 3 UKR      1992 76196   NA    NA Ukraine
## 4 UKR      1993 75896   NA    NA Ukraine
## 5 UKR      1994 70882   NA    NA Ukraine
## 6 UKR      1995 63752   NA    NA Ukraine
## 7 UKR      1996 59080   NA    NA Ukraine
## 8 UKR      1997 54433   NA    NA Ukraine
```

```
## 9 UKR      1998 49938    NA    NA Ukraine
## 10 UKR     1999 47600    NA    NA Ukraine
## # ... with 21 more rows
```

You should see that the information for Ukraine starts in 1990 in the `data_pt` datafile but only in 1994 in the `data_acc` datafile. You can also see that the Ukraine only has information on rail travel, but not road travel.

You could decide, as we are merging data, to only keep information which is available from both datafiles, or to keep all information (Year-country combinations) which are available in at least one of the datasets. Here we decide that we want to do the latter, i.e. keep all information for now, even if it cannot be matched from the other file. So in the above case that means that we do want the information from 1990 to 1993 from Ukraine included. In effect that means that our new dataset, `data_merge` should have at least 2886 rows.

Which of the following commands does that? (Note that, as the information on the basis of which we match, country and year, is saved in identically named columns in both datafiles, we do not have to use the `by.x` and `by.y` options in the merge function).

```
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = TRUE)
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = FALSE)
data_merge <- merge(data_pt,data_acc, all.x = FALSE, all.y = FALSE)
data_merge <- merge(data_pt,data_acc, all.x = FALSE, all.y = TRUE)
```

```
data_merge <- merge(data_pt,data_acc, all.x = TRUE, all.y = TRUE)
```

If you have done this correctly, the new datafile should have 2937 observations and 7 variables.

Task 5: Calculate country averages

We now have multiple years of data for the countries in our dataset, `data_avg`. The next step is to calculate the averages for our four substantial variables (Deaths_p1M, RAIL, ROAD and INLAND).

```
data_avg <- XXXX %>% XXXX(country,LOCATION) %>%
  XXXX(avg_rail = mean(XXXX,na.rm = XXXX),
        avg_road = XXXX,
        avg_inland = XXXX,
        avg_deaths_p1M = XXXX)

data_avg <- data_merge %>% group_by(country,LOCATION) %>%
  summarise(avg_rail = mean(RAIL,na.rm = TRUE),
            avg_road = mean(ROAD,na.rm = TRUE),
            avg_inland = mean(INLAND,na.rm = TRUE),
            avg_deaths_p1M = mean(Deaths_p1M,na.rm = TRUE))
```

You got it right, if your result, `data_avg` has 60 rows and you can replicate the following information:

```
head(data_avg,10)
```

```
## # A tibble: 10 x 6
## # Groups:   country [10]
##   country LOCATION avg_rail avg_road avg_inland avg_deaths_p1M
##   <chr>      <chr>      <dbl>    <dbl>    <dbl>    <dbl>
## 1 Albania   ALB         242.    6075.    6177.     9.86
## 2 Argentina ARG        7765.   34058.   49191.    12.7
## 3 Armenia   ARM         44.8    1945.    1987.    10.1
## 4 Australia AUS       11825.  224331.  236156.     6.68
## 5 Austria   AUT        8785.   55652.   62935.     8.81
## 6 Azerbaijan AZE        1281.   12963.   13920.     9.72
## 7 Belarus   BLR       11235.     NaN     NaN      13.4
```

##	8	Belgium	BEL	8161.	98049.	106104.	9.51
##	9	Bosnia & Herzegovina	BIH	691.	NaN	NaN	8.94
##	10	Bulgaria	BGR	4680.	24248.	31212.	11.2

As you can see, not all countries will have information on all variables.

As mentioned above, the information coming from the passenger transport file are not standardised by population size (`avg_rail`, `avg_road` and `avg_inland`), while `avg_deaths_p1M` already is. We now want to standardise the three variables coming from the passenger transport file by population size. For that we need to import a file with population information. That information (for more than 200 countries) is in `CountryInfo.csv`. Import that file and merge all the available country info into `data_avg` only keeping the 60 rows which we have in `data_avg`.

Note that you should match by the three letter country code which in `CountryInfo.csv` is labelled as `countryCode`. You will therefore, now have to use the `by.x` and `by.y` options in the `merge` function.

```
data_countries = read_csv("XXXX", na = "XXXX" )
data_avg <- merge(data_avg, XXXX,
                  by.x = "XXXX", by.y = "XXXX",
                  all.x = XXXX, all.y = XXXX)
```

```
data_countries = read_csv("CountryInfo.csv", na = "NA" )
data_avg <- merge(data_avg, data_countries,
                  by.x = "LOCATION", by.y = "countryCode",
                  all.x = TRUE, all.y = FALSE)
```

Now check the names of the variables in your datafile.

```
names(data_avg)
```

##	[1]	"LOCATION"	"country"	"avg_rail"	"avg_road"
##	[5]	"avg_inland"	"avg_deaths_p1M"	"popData2019"	"continentExp"
##	[9]	"Land_Area_sqkm"	"HealthExp"	"GDPpc"	"Obese_Pcent"
##	[13]	"Over_65s"	"Diabetis"		

You should have 14 variables. Also confirm that you do have 60 rows of data. All the new country data are from the year 2019.

Task 6: Standardisation

We now need to calculate the standardised personal transport variables.

```
data_avg <- data_avg %>% XXXX(
  avg_rail_pp = avg_rail/popData2019 * 1000000,
  avg_road_pp = XXXX,
  avg_inland_pp = XXXX)

data_avg <- data_avg %>% mutate(
  avg_rail_pp = avg_rail/popData2019 * 1000000,
  avg_road_pp = avg_road/popData2019 * 1000000,
  avg_inland_pp = avg_inland/popData2019 * 1000000)
```

You get it right if you can replicate the following.

```
data_avg %>% filter(LOCATION %in% c("FRA", "ESP")) %>%
  select(LOCATION, avg_rail_pp, avg_road_pp, avg_inland_pp)

## LOCATION avg_rail_pp avg_road_pp avg_inland_pp
```

```
## 1      ESP      408.7568      5487.916      5897.644
## 2      FRA     1119.4566      9651.300     10770.756
```

So the average number of miles traveled by rail per person in Spain is 409 miles per year. The average number of miles traveled on road per person in France is 9651. Convince yourself that the above calculation ($\text{avg_rail_pp} = \text{avg_rail}/\text{popData2019} * 1000000$) did deliver miles per person. Recall that the travel distance coming from the passenger transport datafile was measured in Millions of miles. The population is the actual population. If we didn't multiply by 1,000,000 we would get a measure of average millions of miles traveled per person.

Lastly, we will rescale the `GDPpc` variable. It is currently defined in USD, but to simplify later analysis we want to express it in 1000 USD. So for instance Albania's `GDPpc` is USD 5224, but we want to express it as 5.224 [1000 USD].

```
data_avg <- data_avg %>% mutate(GDPpc = GDPpc/1000)
```

Look at the dataset to confirm that your operation was successful.

Task 7: Creating League Tables

Create a league table of the countries in your dataset where people travel most and least by rail and road. You want to display the 10 countries with the most and the 10 countries with the least average travel per person in both categories.

Here is the code for the table for the top 10 rail travel nations.

```
task_7a <- data_avg %>%
select(country, avg_rail_pp) %>% arrange(desc(avg_rail_pp))
head(task_7a,10)
```

```
##      country avg_rail_pp
## 1      Japan  2908.5449
## 2 Switzerland  1591.2104
## 3      Latvia  1348.2763
## 4      Russia  1318.4358
## 5      Belarus  1188.9556
## 6      Ukraine  1153.9672
## 7      France  1119.4566
## 8      Hungary  1065.6674
## 9      Austria   986.9251
## 10 Kazakhstan   915.7539
```

Repeat this for the Top 10 road travel nations and equally find similar tables for the nations travelling least (but non-zero amounts) by rail and road.

```
task_7b <- data_avg %>%
  select(country, avg_road_pp) %>% arrange(desc(avg_road_pp))
head(task_7b,10)
```

```
country avg_road_pp
```

```
1 United States 14614.973 2 Iceland 13806.988 3 Canada 13188.024 4 Lithuania 11804.369 5 Slovenia 11211.417
6 Finland 10526.496 7 Denmark 10459.718 8 Italy 10183.750 9 New Zealand 9841.385 10 Sweden 9709.163
```

```
task_7c <- data_avg %>%
  select(country, avg_rail_pp) %>% arrange(avg_rail_pp)
head(task_7c,10)
```

```
country avg_rail_pp
```


1 Armenia 15.12183 2 Mexico 21.53597 3 Chile 35.75820 4 Canada 36.60707 5 Turkey 76.57295 6 United States 79.73147 7 Albania 84.87356 8 Uzbekistan 93.37132 9 Morocco 95.22897 10 North Macedonia 104.12931

```
task_7d <- data_avg %>%
  select(country, avg_road_pp) %>% arrange(avg_road_pp)
head(task_7d,10)
```

```
country avg_road_pp
```

1 Montenegro 174.8040 2 Armenia 656.4899 3 Argentina 753.5608 4 Russia 1076.1260 5 Azerbaijan 1278.4664
6 Georgia 1434.4690 7 Uzbekistan 1701.8264 8 Turkey 1964.1749 9 Albania 2134.6885 10 Greece 2660.8912

Task 8: Estimate regression models - Version 1

We shall estimate the following three regression models (`mod1`)

$$avg_deaths_p1M = \gamma_0 + \gamma_1 GDPpc + w$$

and (`mod2`)

$$avg_deaths_p1M = \alpha_0 + \alpha_1 avg_rail_pp + v$$

and (`mod3`)

$$avg_deaths_p1M = \beta_0 + \beta_1 avg_road_pp + \beta_2 GDPpc + u$$

```
mod1 <- lm(XXXX ~ XXXX, data = data_avg)
mod2 <- lm(XXXX)
mod3 <- lm(XXXX)
stargazer_HC(mod1,mod2,mod3,omit.stat = "f")
```

```
mod1 <- lm(avg_deaths_p1M ~ GDPpc, data = data_avg)
mod2 <- lm(avg_deaths_p1M ~ avg_rail_pp, data = data_avg)
mod3 <- lm(avg_deaths_p1M ~ avg_road_pp+GDPpc, data = data_avg)
stargazer_HC(mod1,mod2,mod3,type_out="latex",omit.stat = "f")
```

% Table created by stargazer v.5.2.3 by Marek Hlavac, Social Policy Institute. E-mail: marek.hlavac at gmail.com % Date and time: Thu, Mar 02, 2023 - 08:41:56

If you have done this correctly, you will find that that your estimated constant for `mod1` is 11.777, the estimated constant for `mod2` is 9.911 and that `mod3` is estimated with 44 observations.

Task 9: Plot the data

Now we want to plot the `GDPpc` versus the `avg_deaths_p1M` data in a scatter plot. Replicate the graph below.

You should change the Axis labels and add a title as shown below. If you google you should find the appropriate commands to do so. (Think carefully about the search terms.) Also, what does the `geom_smooth(method='lm')` line achieve? Figure that out by running the code without that line.

```
ggplot(data_avg, aes(x=XXXX,y=XXXX)) +
  geom_point(color = "blue") +
  geom_smooth(method='lm') +
  XXXX +
```

Table 1:


	<i>Dependent variable:</i>		
	avg_deaths_p1M		
	(1)	(2)	(3)
avg_road_pp			0.0003 (0.0002)
GDPpc	-0.066*** (0.020)		-0.119*** (0.024)
avg_rail_pp		0.0001 (0.001)	
Constant	11.777*** (0.697)	9.911*** (0.663)	11.328*** (1.097)
Observations	58	55	44
R ²	0.235	0.0002	0.342
Adjusted R ²	0.221	-0.019	0.310
Residual Std. Error	3.096 (df = 56)	3.464 (df = 53)	2.952 (df = 41)

Note:


*p<0.1; **p<0.05; ***p<0.01
Robust standard errors in parenthesis

XXXX +
XXXX

```
ggplot(data_avg, aes(x=GDPpc,y=avg_deaths_p1M)) +
  geom_point(color = "blue") +
  geom_smooth(method='lm') +
  ggtitle("Accidents v GDPpc") +
  ylab("Accidents per Million population") +
  xlab("GDP per capita")
```



Finally another searching challenge for you. The picture below has added data labels to the above plot. Find a way to achieve this. It is not important that the graph looks exactly like the one below, but you do want to find out how to add data labels. Dr. Google is your friend!



And does anyone know what is going on in Luxembourg?

END OF INSTRUCTIONS